

PROFESSIONAL RESEARCH ANALYSIS

"Evaluation of LLMs on Bloom's Taxonomy Classification"

Senior Researcher Assessment & Strategic Guidance

Author: Muhammad Daniyal

Date: April 2026

Analyst Role: Senior Research Advisor

EXECUTIVE SUMMARY

This is a **methodologically sound but strategically mid-tier positioning** of an LLM benchmarking study on Bloom's Taxonomy classification. The work demonstrates:

 Strengths:

- First published QLoRA fine-tuning on Monash CLO dataset
- Broadest LLM sweep (13 models, zero/few-shot + fine-tuning)
- Rigorous experimental design with proper metrics
- Valuable Gemma-2 failure diagnosis (reproducibility artifact)
- Strong technical execution (93.22% accuracy, 0.9206 macro F1)

 Critical Gaps:

- Performance is 3 points below BERT SOTA (Li/Sridhar 2022: F1 0.95)
- No explainability component (immediate competitive disadvantage vs Sebbaq 2023)
- No educator validation (fails pedagogical rigor test)
- Single-label only (dropped 12% multi-label cases)
- Narrative weak on novelty (15 papers from 2020-2026 all do similar benchmarking)

Verdict: Publishable but incremental without additions. The work needs 2-3 "teeth" to be competitive at CAEAI or IEEE TLT. The roadmap provided in slides 39-40 is sound; following it will elevate this to a strong contribution.

WHAT'S WORKING WELL

1. Experimental Rigor (Slides 7-12)

Dimension	Assessment	Grade
Dataset handling	Correct stratified 85/15, multi-label filtering explicit	A
Metrics	9 metrics including MCC, Cohen's κ, balanced accuracy	A-
Baseline establishment	Classical ML floor at 72% is reasonable comparator	A
Model coverage	13 open-weight LLMs across three families (Mistral, Qwen, Gemma)	A
Fine-tuning config	Identical hyperparameters across all three models isolates architecture	A
Reproducibility artifacts	Prompt templates, training curves, loss data all documented	B+

Judgment: The empirical backbone is solid. This work could be reproduced faithfully from the slides.

2. Technical Quality of Results (Slides 19-27)

- **Mistral-7B winner:** 93.22% accuracy, 0.9206 macro F1 across 6 classes, F1 range 0.86–0.95 per class
- **Qwen2.5-7B strong second:** 92.47%, within 0.8 points of winner
- **Per-class balance:** Even rare "Remember" class (133 test samples) achieves F1 0.86 — excellent imbalance handling
- **Confusion matrix quality:** No pathological patterns

Judgment: These numbers are credible, well-distributed across classes, and genuinely competitive with traditional ML (72% LogReg baseline).

3. Failure Analysis (Slide 25)

The Gemma-2 diagnosis is valuable:

Symptom: Precision ≈ 1.0, Recall ≈ 0.06 — only labels 6% of higher-order outcomes

Root causes identified:

1. SDPA/FlashAttention incompatible with soft-capping
2. Chat template doesn't append `<eos>`
3. `pad_token = eos_token` masks stop token in loss
4. Learning rate 2e-4 too aggressive for soft-capped architecture
5. fp16 overflows near soft-cap

5-step fix documented: `attn_implementation="eager"`, distinct pad token, explicit EOS, `LR→5e-5`, `bf16=True`

Why this matters: This is a real problem that community members hit repeatedly (HF issue #32110, axolotl #1921). Documenting before/after with metrics is a reproducibility artifact worth publishing. **Suitable for IEEE TLT or a workshop paper on training recipes.**

CRITICAL GAPS

⚠ Gap 1: NO EXPLAINABILITY (High Priority)

Current state: Zero mention of SHAP, attention visualization, or interpretability.

Competitive landscape:

- **Sebbaq & El Faddouli (2023):** MTBERT + LIME, 97.7% accuracy — **had XAI**
- **Huber & Niklaus (2025, COLING):** GPT-4/Claude-3 benchmarks, discussed what makes Create level hard
- **Almatrafi & Johri (2025, CAEAI):** Alluded to CAEAI's preference for interpretable classifiers

Reviewer expectation (2025-2026): Especially for an education-focused venue, explaining *why* the model predicts "Evaluate" for a CLO is as important as getting it right.

Required for competitiveness:

- SHAP (TransSHAP for text) or Integrated Gradients
- Attention head analysis (which attention patterns align with Bloom levels?)
- **Verb-Level Alignment Score:** measure if attention flows to the action verb
- ERASER faithfulness metrics (sufficiency, comprehensiveness)

Effort: 2-3 weeks to add (SHAP + seaborn visualizations).

⚠ Gap 2: NO EDUCATOR VALIDATION (Critical)

Missing: Pedagogical expert review of predictions and explanations.

Benchmark (Sebbaq 2023): No educator panel, but LIME explanations were shown to be interpretable.

Your gap: Even without formal panel, adding 3-5 instructional-design experts rating 30-60 predictions would:

- Compute Fleiss' kappa for inter-rater agreement
- Ask if explanations match pedagogical intuition
- Validate that "verb prominence" is the right signal

Educational measurement standard: Krippendorff $\alpha = 0.67$ for Bloom labeling (Marzi 2024). Your model at 93% is *above human agreement*, but educators need to validate the reasoning.

Effort: 3-4 weeks (recruiting, annotation, analysis).

⚠ Gap 3: SINGLE-LABEL ONLY (Moderate Priority)

Issue: Dropped ~2,600 multi-label CLOs (12% of data).

Why it matters:

- **Pedagogical truth:** Many CLOs legitimately target multiple levels. "Evaluate the design of X and propose improvements" = Evaluate + Create.
- **Co-occurrences in your data:**
 - Evaluate ↔ Analyze: 522 examples
 - Evaluate ↔ Apply: 399
 - Create ↔ Evaluate: 376

These are the *most interesting* examples for Bloom measurement — exactly where human annotation is noisiest.

Competitive disadvantage:

- Li/Sridhar (2022) *built multi-label classifiers* (6 binary heads).
- Regression from their approach weakens your paper's scope.

Required for competitiveness:

- Restore multi-label cases
- One-vs-rest heads with per-class calibration (`sklearn.calibration.CalibratedClassifierCV`)
- Report subset accuracy, Hamming loss, per-co-occurrence confusion

Effort: 2-3 weeks.

⚠ Gap 4: NARRATIVE WEAKNESS ON NOVELTY

The "many-model bake-off" is exhausted in literature:

- Mohammed & Omar (2020) — TF-IDF, ML models
- Kumar et al. (2025) — SVM + augmentation on 600 samples
- Alshanqiti/Alammary (2025) — DistilBERT ensembles
- Sebbaq (2023) — MTBERT with LIME
- **You (2026)** — 13 open-weight LLMs, zero/few-shot/QLoRA

The reviewer will ask: **"What do the 13 models teach us that the prior 15 papers did not?"**

Current answer from your presentation:

-  First QLoRA on Monash CLO (novel methodology)
-  Broadest LLM sweep (novel scope)
-  Results are 3 pts below BERT SOTA (not novel achievement)
-  No explainability added (not novel insight)

Risk: "Incremental benchmarking of an existing dataset with new architectures" is vulnerable to desk rejection at CAEAI or IEEE TLT.

Solution: Add explainability + multi-label + cross-dataset analysis. That becomes a *new research agenda*, not a "model comparison."

 Gap 5: COMPARISON STUDY POSITIONING UNCLEAR

Direct competitor: Almatrafi & Johri (2025, CAEAI)

Dimension	Almatrafi 2025	This Work
Approach	GPT-4 (closed, SOTA LLM)	Qwen2.5-7B + Mistral (open-weight)
Dataset	~1,000 CLOs	18,773 CLOs
XAI	Not mentioned	None (gap)
Multi-label	Not clear	Single-label only (gap)
Reproducible	No (closed API)	Yes (HF models)

Reviewer perspective: "Almatrafi already proved GPT-4 works. Why should I accept a 93% open-weight model on a larger dataset if you don't explain *how* it works?"

Answer needed: Your value is reproducibility + explainability + scale, not raw accuracy.

POSITIONING ANALYSIS: WHERE YOU STAND

State-of-the-Art Timeline

2020-2022: BERT Era

- └ Li/Sridhar (2022, EDM): BERT-base, F1 0.95 ← MONASH SOTA

2023-2024: BERT + XAI / Small LLM Evals

- └ Sebbag (2023): MTBERT + LIME, 97.7%
- └ Chindukuri (2024): Cross-dataset transfer, 78.44% OOD

2025-2026: LLM Era (Current)

- └ Almatrafi (2025, CAEAI): GPT-4 prompting
- └ Kumar (2025, arXiv): SVM + aug, 94%
- └ Alammary (2025, MDPI): DistilBERT, 96% (small data)
- └ YOU (2026): QLoRA Mistral, 93.22% ← HERE

Your position: In the middle of a crowded field. Not below (93% > 72% baseline), not above (93% < 95-97% published SOTA).

To move up, choose one:

Path	Effort	Risk	Payoff	Target Venue
A: Explainability + Educator Validation	4-5 weeks	Low	High (unique)	CAEAI
B: Multi-Label + Cross-Dataset	4-5 weeks	Low	High (novel framing)	IEEE TLT
C: Gemma-2 Recipe Publication	1-2 weeks	Low	Medium (niche)	TLT Workshop
D: Benchmark Suite + Leaderboard	6-8 weeks	Medium	Medium (infrastructure)	CAEAI special section

SPECIFIC RECOMMENDATIONS FOR NEXT WORK

IMMEDIATE (Next 2 Weeks)

1. Restore Multi-Label Cases

```
python
```

```
# Current: 18,773 single-label samples
```

```
# Target: 21,379 all-label samples
```

```
# Approach:
```

```
# - One-vs-rest heads: 6 independent binary classifiers
```

```
# - BCE with logits loss
```

```
# - Threshold each head at 0.5 (or tune per-class via calibration)
```

```
# Metrics:
```

```
# - Subset accuracy (all 6 match exactly): ~30-40% expected
```

```
# - Hamming loss: avg # wrong bits per sample
```

```
# - Per-label F1
```

```
# - Per-co-occurrence confusion: especially Evaluate-Analyze, Create-Evaluate
```

```
# Why: Restores pedagogical realism + differentiates from single-label papers
```

2. Acquire Cross-Dataset Samples

```
- Yahya et al. (2017): 600 LOs, 4-level Bloom (English)
```

```
- Mohammed & Omar (2020): 141 CLOs, 6-level
```

```
- NCERT (India): ~3,000 text excerpts → extract LOs
```

```
- AAFAQ (Arabic): 5,009 questions, 4-level
```

```
Rationale: Leave-one-dataset-out evaluation measures generalization
```

WEEK 2-4: EXPLAINABILITY FOUNDATION

Build SHAP + Attention Analysis (Top Priority)

```
python
```

```
# Stage 1: SHAP TransSHAP for text
from shap import Explainer
explainer = Explainer(model, tokenizer, device='cuda')
shap_values = explainer(clo_texts[:100]) # Sample 100 CLOs

# Key metrics:
# - Feature importance by Bloom level
# - Does the model attend to action verbs (define, analyze, create)?
# - Saliency maps per class

# Stage 2: Attention Head Analysis
# Extract attention weights from final layer
# Compute attention flow: [CLS] → VERBS → [END]
# Measure: Attention(verb) / Attention(all_tokens)

# Stage 3: ERASER Faithfulness
# Sufficiency: remove top-K SHAP tokens, does F1 drop?
# Comprehensiveness: keep top-K SHAP tokens, does F1 stay high?

# Output: One figure per Bloom level showing:
# - Top 10 SHAP features
# - Attention patterns
# - Faithfulness trade-off curve
```

Why this matters:

- SHAP alone will differentiate you from Almatrafi (2025, no XAI)
- Educators immediately understand token-level salience
- Venues (CAEAI, IEEE TLT) heavily weight interpretability in 2025

WEEK 4-6: EDUCATOR VALIDATION PIPELINE

Recruit & Run Panel (3-5 experts)

Target recruitment:

- 1-2 curriculum designers
- 2-3 subject-matter educators (science, engineering, liberal arts mix)
- Budget: \$300-500 total (gift cards for time)

Study design:

- Show model prediction + top-3 SHAP-highlighted words
- Ask:
 1. Does this Bloom level match the CLO? (yes/no/unsure)
 2. If wrong, what level should it be?
 3. How clear is the explanation? (1-5 Likert)
- Sample: 60-120 CLOs (30 from each worst-performing class)

Analysis:

- Fleiss' kappa (multi-rater agreement)
- Cohen's weighted kappa (with distance between levels)
- ICC(3,k) intra-class correlation
- Category agreement matrix

Expected outcome:

- $\kappa \approx 0.70-0.85$ (your model exhibits high inter-rater agreement, approaching human expert consensus of $\kappa \approx 0.67$ from Marzi 2024)

This is publishable as:

"Educators agree 78% of the time with the model on Bloom level assignments, and when they disagree, model explanations help educators understand the reasoning."

WEEK 6-8: CROSS-DATASET + CROSS-TAXONOMY

Leave-One-Dataset-Out Evaluation

python

```
datasets = {
    'monash_clo': 18773,
    'yahya_2017': 600,
    'mohammed_omar': 141,
    'ncert_extracted': 3000,
    'aafaq_arabic': 5009
}

for test_dataset in datasets:
    train_data = concat([datasets[d] for d in datasets if d != test_dataset])

    # Train QLoRA Mistral-7B on 26,623 samples
    # Test on ~2,000 samples from test_dataset

    # Report:
    # - Accuracy, Macro F1
    # - Per-class F1 (Create is hardest across datasets?)
    # - Error patterns: does Understand→Apply confusion persist across datasets?

# Expected insight:
# - In-domain (Monash): 93.2%
# - OOD (Yahya, small): 78-82%
# - OOD (NCERT English, large): 85-88%
# - OOD (AAFAQ Arabic, new language): ~45-55% (significant drop)
```

Why this matters:

- Generalization is the unsung hero of educational NLP
- Chindukuri (2024) already showed transfer is hard (78% OOD)
- You prove that fine-tuned LLMs + explainability enable robust cross-dataset performance
- **NEW CONTRIBUTION:** No prior paper on fine-tuned LLM generalization for Bloom

WEEK 8-10: CROSS-TAXONOMY ANALYSIS (High Impact)

Add SOLO + DOK Dimensions

The **single biggest opportunity** in educational NLP right now is bridging Bloom, SOLO (Structure of Observed Learning Outcomes), and DOK (Webb's Depth of Knowledge).

Bloom: 6 levels (traditional, cognitive)
SOLO: 5 levels (information complexity: unistructural → extended abstract)
DOK: 4 levels (cognitive demand in assessment)

Joint modeling:

- Annotate ~500 CLOs with all three taxonomies
- Build Bloom-SOLO-DOK relationship graph
- Fine-tune model to predict all three jointly

Insight: Different taxonomies measure different things

- Bloom ≈ action verb (analyze, create)
- SOLO ≈ complexity of thinking (surface vs deep)
- DOK ≈ cognitive demand in assessment (recall vs reasoning)

A CLO could be:

- Bloom: Remember (low)
- SOLO: Multistructural (medium) – multiple disconnected facts
- DOK: Level 3 (high) – requires strategic thinking to connect facts

Why publish this:

- DOK is almost **entirely absent** from NLP literature (except You et al. 2025 on APR)
- SOLO is barely explored (only Biggs/Collis papers are old)
- This is a **virgin territory** — highest novelty potential

14-WEEK EXPERIMENTAL ROADMAP (Detailed)

Week	Task	Deliverable	Owner
1-2	Data consolidation: restore multi-label, acquire 4 new datasets	<code>data_consolidated.csv</code> (21.3K + 8.7K = 30K)	You
2-3	Multi-label modeling: one-vs-rest heads, calibration	<code>mistral_multilabel.py</code> , metrics table	You
3-5	SHAP + attention analysis: TransSHAP, attn heads, ERASER	<code>shap_analysis.pdf</code> (6 pages, 1 fig per Bloom level)	You
5-6	Gemma-2 fix: test with corrected config, before/after metrics	<code>gemma_fix_report.md</code> , TLT paper draft	You
6-8	Educator validation: recruit panel, run annotation, κ analysis	<code>educator_validation.csv</code> , <code>agreement_report.pdf</code>	You + 3-5 educators
8-10	Cross-dataset LODO: train on 26K, test each on OOD	<code>cross_dataset_results.csv</code> , transfer loss curves	You
8-10	Cross-taxonomy: SOLO/DOK annotation on 500 samples, joint model	<code>joint_bloom_solo_dok.json</code> , confusion matrices	You (or subcontract annotation)
10-12	Writing: draft CAEAI paper (explainability focus)	<code>paper_v1_caeai.pdf</code> (10-12 pages)	You
11-12	Reviewer prep: re-bench vs Almatrafi (2025), identify unique angles	<code>novelty_audit.md</code> , response to anticipated reviewer comments	You
12-14	Final revisions, TLT workshop submission (Gemma-2 recipe)	<code>paper_v2_tlt_recipe.pdf</code> (5-6 pages)	You

PUBLICATION STRATEGY

Primary Target: CAEAI 2026 (Fall Deadline)

Title suggestion:

"Explaining Bloom's Taxonomy Classifications with Fine-Tuned LLMs and Educator Validation: A QLoRA Study on 30K Course Learning Outcomes"

Core narrative:

1. **Problem:** GPT-4 is closed. BERT models from 2022 don't explain their decisions. Educators need transparent, reproducible, interpretable classifiers.
2. **Solution:** Fine-tuned open-weight LLMs + SHAP explanations + educator validation.
3. **Evidence:**
 - Mistral-7B at 93.2% (within 3 pts of BERT SOTA)
 - SHAP top-3 tokens match pedagogical intuition 78% of time
 - 5 educators agree on 81% of predictions
4. **Scope:** 30K CLOs across 5 datasets, cross-dataset + cross-taxonomy
5. **Impact:** First explainable fine-tuned LLM Bloom classifier, reproducible on consumer GPUs

Why CAEAI accepts:

- 2025 papers (Almatrafi, others) proved LLMs work for Bloom
- 2026 papers need to answer: how do we *trust* them?
- Your explainability + validation directly answer this

Estimated acceptance: 65-75% (high novelty + rigor, but not groundbreaking performance)

Secondary Target: IEEE TLT 2026

Option A: Full technical paper (parallel to CAEAI)

"Multi-Label Bloom's Taxonomy Classification with Cross-Dataset Generalization: A QLoRA Fine-Tuning Study"

Focus: Engineering + generalization, not explainability.

Option B: Workshop paper on Gemma-2 recipes

"Training Gemma-2-9B-Instruct for Educational Text Classification: A Reproducibility Study and Configuration Guide"

- Short (4-6 pages)
 - Practical value for practitioners
 - Suitable for IEEE TLT workshop track
 - Lower barrier to acceptance
-

RISK ASSESSMENT & MITIGATION

Risk	Probability	Impact	Mitigation
CAEAI rejects as incremental	30%	High	(1) Explainability is core, (2) educator validation is novel, (3) position as XAI paper not benchmarking
Educator panel gives poor κ	20%	Medium	Choose educators carefully; provide SHAP visualizations; budget for 2+ rounds of annotation
Cross-dataset OOD drop is steep	40%	Low	Expected; frame as "domain-specific refinement needed" and show fix (domain adaptation)
AAFAQ Arabic data doesn't generalize	50%	Low	Arabic Bloom is understudied; any result is publishable; frame as "cross-lingual challenge"
Almatrafi 2025 (CAEAI) cites you first	5%	Medium	Impossible to control; emphasize your differences (open vs closed, explainability, validation)

FINAL RECOMMENDATION

What to Do Right Now (Next 2 Weeks)

- Restore multi-label cases** (2 days)
 - This removes a major objection from reviewers
 - Shows you're serious about pedagogical realism
- Run SHAP analysis on Mistral** (1 week)
 - Pick 30 representative CLOs (5 per Bloom level)
 - Generate SHAP feature importance + attention heatmaps
 - Write a brief (2-page) memo: "Are Bloom classifiers discovering action verbs?"
 - Share with 1-2 educators for quick feedback
- Draft Gemma-2 TLT paper** (1 week)
 - Before/after config, metrics, reproducibility notes
 - Not a high-impact paper, but solid engineering contribution
 - De-risk publication: IEEE TLT has high acceptance for recipes + diagnostics
- Recruit educator panel** (in parallel, weeks 2-4)
 - Start conversations now; don't wait

Strategic Narrative Going Forward

Don't lead with: "We fine-tuned 13 LLMs and got 93%."
Lead with: "We built the first open, explainable, validated Bloom classifier. Here's why educators should trust it."

This shifts you from **benchmarking** (commoditized, many papers) to **trustworthy AI for education** (urgent, few papers, high impact).

CONCLUSION

Your work is **technically solid and well-executed**, but needs **narrative and novelty amplification** to be competitive at top venues. The roadmap above (14 weeks, 4-5 major additions) takes you from "solid benchmarking paper" to "landmark study in explainable educational AI."

Key insight: The barrier isn't accuracy (93% is good). The barrier is answering: *Why should educators use an AI system that I can't explain?*

That question, once answered rigorously, makes this paper publishable and impactful.

Next Step: Schedule a working session to confirm:

1. ☒ Multi-label restoration (priority: high)
2. ☒ SHAP analysis (priority: high)
3. ☒ Educator recruitment contact list (priority: medium)
4. ☒ Cross-dataset acquisition plan (priority: medium)
5. ☒ Timeline adjustment based on available effort

Good luck with the next phase. This is solid research with a clear path to publication.